

# Contents

[Introduction](#)

[Prerequisites](#)

[Requirements](#)

[Components Used](#)

[Conventions](#)

[Load Sharing with the Loopback Address as a BGP Neighbor](#)

[Network Diagram](#)

[Configurations](#)

[Verify](#)

[Troubleshoot](#)

[Load Sharing When Dual-Homed to One Internet Service Provider \(ISP\) Through a Single Local Router](#)

[Network Diagram](#)

[Configurations](#)

[Verify](#)

[Troubleshoot](#)

[Load Sharing When Dual-Homed to One ISP Through Multiple Local Routers](#)

[Network Diagram](#)

[Configurations](#)

[Verify](#)

[Verification When Both Links Between AS 11 and AS 10 Are Up](#)

[Verification When the R101-R103 Link Fails](#)

[Troubleshoot](#)

[Load Sharing When Multihomed to Two ISPs Through a Single Local Router](#)

[Network Diagram](#)

[Configurations](#)

[Verify](#)

[Troubleshoot](#)

[Load Sharing When Multihomed to Two ISPs Through Multiple Local Routers](#)

[Network Diagram](#)

[Configurations](#)

[Verify](#)

[Troubleshoot](#)

[Related Information](#)

## Introduction

Load sharing allows a router to distribute the outgoing and incoming traffic among multiple paths. The paths are derived either statically or with dynamic protocols, such as:

- Routing Information Protocol (RIP)
- Enhanced Interior Gateway Routing Protocol (EIGRP)

- Open Shortest Path First (OSPF) Protocol
- Interior Gateway Routing Protocol (IGRP)

By default, Border Gateway Protocol (BGP) selects only a single best path and does not perform load balancing. This document shows how to perform load sharing in different scenarios with the use of BGP. For additional information about load balancing, refer to [How Does Load Balancing Work?](#).

## Prerequisites

### Requirements

Ensure that you meet these requirements before you attempt this configuration:

- Knowledge of [BGP Best Path Selection Algorithm](#)
- Knowledge of [Configuring BGP](#)

### Components Used

This document is not restricted to specific software and hardware versions.

The information in this document was created from the devices in a specific lab environment. All of the devices used in this document started with a cleared (default) configuration. If your network is live, make sure that you understand the potential impact of any command.

### Conventions


Refer to [Cisco Technical Tips Conventions](#) for more information on document conventions.

## Load Sharing with the Loopback Address as a BGP Neighbor

This scenario shows how to achieve load sharing when there are multiple (up to a maximum of six), equal-cost links. The links are terminated in one router at a local autonomous system (AS) and in another router at a remote AS in a single-homed BGP environment. The [Network Diagram](#) serves as an example.

### Network Diagram

This section uses this network setup:



## Configurations

This section uses these configurations:

- [RouterA](#)
- [RouterB](#)

### RouterA

```
interface loopback 0
 ip address 1.1.1.1 255.255.255.0

interface serial 0
 ip address 160.20.20.1 255.255.255.0
 no ip route-cache

interface serial 1
 ip address 150.10.10.1 255.255.255.0
 no ip route-cache

router bgp 11
 neighbor 2.2.2.2 remote-as 10
 neighbor 2.2.2.2 update-source loopback 0
 !--- Use the IP address of the loopback interface for TCP connections. neighbor 2.2.2.2 ebgp-multihop
 !--- You must configure ebgp-multihop whenever the external BGP (eBGP)
 !--- connections are not on the same network address. router eigrp 12 network 1.0.0.0 network 150.10.0.
 network 160.20.0.0 no auto-summary
```

### RouterB

```
interface loopback 0
 ip address 2.2.2.2 255.255.255.0

interface serial 0
 ip address 160.20.20.2 255.255.255.0
 no ip route-cache

interface serial 1
 ip address 150.10.10.2 255.255.255.0
 no ip route-cache

router bgp 10
 neighbor 1.1.1.1 remote-as 11
 neighbor 1.1.1.1 update-source loopback 0
 !--- Use the IP address of the loopback interface for TCP connections. neighbor 1.1.1.1 ebgp-multihop
 !--- You must configure ebgp-multihop whenever the eBGP connections
 !--- are not on the same network address. router eigrp 12 network 2.0.0.0 network 150.10.0.0 network
 160.20.0.0 no auto-summary
```

**Note:** You can use static routes in the place of a routing protocol in order to introduce two equal-cost paths to reach the destination. In this case, the routing protocol is EIGRP.

## Verify

Use this section to confirm that your configuration works properly.

The [Cisco CLI Analyzer](#) ([registered](#) customers only) supports certain **show** commands. Use the Cisco CLI Analyzer to view an analysis of **show** command output.

The output of the **show ip route** command shows that both of the paths to the 2.2.2.0 network are learned via EIGRP. The output of the **traceroute** command indicates that the load is distributed between two serial links. In this scenario, load sharing occurs on a per-packet basis. You can issue the **ip route-cache** command on the serial interfaces to do load sharing on a per-destination basis. You can also configure per-packet and per-destination load balancing with Cisco Express Forwarding. For more information on how to configure Cisco Express Forwarding, refer to [Configuring Cisco Express Forwarding](#).

```
RouterA# show ip route
!--- Output suppressed. Gateway of last resort is not set 1.0.0.0/24 is subnetted, 1 subnets C
1.1.1.0 is directly connected, Loopback0 2.0.0.0/24 is subnetted, 1 subnets D      2.2.2.0
[90/2297856] via 150.10.10.2, 00:00:45, Serial1
      [90/2297856] via 160.20.20.2, 00:00:45, Serial0
160.20.0.0/24 is subnetted, 1 subnets
C      160.20.20.0 is directly connected, Serial0
150.10.0.0/24 is subnetted, 1 subnets
C      150.10.10.0 is directly connected, Serial1
```

```
RouterA# traceroute 2.2.2.2
```

```
Type escape sequence to abort.
Tracing the route to 2.2.2.2
```

```
1 160.20.20.2 16 msec
   150.10.10.2 8 msec *
```

## Troubleshoot

There is currently no specific troubleshooting information available for this configuration.


## Load Sharing When Dual-Homed to One Internet Service Provider (ISP) Through a Single Local Router

This scenario shows how to achieve load sharing when multiple links exist between a remote AS and a local AS. These links are terminated in one router at the local AS and on multiple routers at remote ASs in a single-homed BGP environment. The [Network Diagram](#) is an example of such a network.

This sample configuration uses the **maximum-paths** command. By default, BGP chooses one best path among the possible equal-cost paths that are learned from one AS. However, you can change the maximum number of parallel equal-cost paths that are allowed. In order to make this change, include the **maximum-paths paths** command under the BGP configuration. Use a number between 1 and 6 for the *paths* argument.

## Network Diagram

This section uses this network setup:



## Configurations

This section uses these configurations:

- [RouterA](#)
- [RouterB](#)
- [RouterC](#)

### RouterA

```

RouterA# show ip route
!--- Output suppressed. Gateway of last resort is not set 1.0.0.0/24 is subnetted, 1 subnets C 1.1.1.0
directly connected, Loopback0 2.0.0.0/24 is subnetted, 1 subnets D      2.2.2.0 [90/2297856] via
150.10.10.2, 00:00:45, Serial1
    [90/2297856] via 160.20.20.2, 00:00:45, Serial0
    160.20.0.0/24 is subnetted, 1 subnets
C      160.20.20.0 is directly connected, Serial0
    150.10.0.0/24 is subnetted, 1 subnets
C      150.10.10.0 is directly connected, Serial1

RouterA# traceroute 2.2.2.2

Type escape sequence to abort.
Tracing the route to 2.2.2.2

 1 160.20.20.2 16 msec
   150.10.10.2 8 msec *
```

### RouterB

```

RouterA# show ip route
!--- Output suppressed. Gateway of last resort is not set 1.0.0.0/24 is subnetted, 1 subnets C 1.1.1.0
directly connected, Loopback0 2.0.0.0/24 is subnetted, 1 subnets D      2.2.2.0 [90/2297856] via
150.10.10.2, 00:00:45, Serial1
    [90/2297856] via 160.20.20.2, 00:00:45, Serial0
    160.20.0.0/24 is subnetted, 1 subnets
C      160.20.20.0 is directly connected, Serial0
    150.10.0.0/24 is subnetted, 1 subnets
C      150.10.10.0 is directly connected, Serial1

RouterA# traceroute 2.2.2.2

Type escape sequence to abort.
Tracing the route to 2.2.2.2

 1 160.20.20.2 16 msec
   150.10.10.2 8 msec *
```

### RouterC

```

RouterA# show ip route
!--- Output suppressed. Gateway of last resort is not set 1.0.0.0/24 is subnetted, 1 subnets C 1.1.1.0
directly connected, Loopback0 2.0.0.0/24 is subnetted, 1 subnets D      2.2.2.0 [90/2297856] via
```

```
150.10.10.2, 00:00:45, Serial1
      [90/2297856] via 160.20.20.2, 00:00:45, Serial0
160.20.0.0/24 is subnetted, 1 subnets
C      160.20.20.0 is directly connected, Serial0
150.10.0.0/24 is subnetted, 1 subnets
C      150.10.10.0 is directly connected, Serial1
```

```
RouterA# traceroute 2.2.2.2
```

```
Type escape sequence to abort.
Tracing the route to 2.2.2.2
```

```
1 160.20.20.2 16 msec
  150.10.10.2 8 msec *
```

## Verify

Use this section to confirm that your configuration works properly.

The [Cisco CLI Analyzer](#) ([registered](#) customers only) supports certain **show** commands. Use the Cisco CLI Analyzer to view an analysis of **show** command output.

The output of the **show ip route** command shows that both of the paths to the 2.2.2.0 network are learned via BGP. The output of the **traceroute** command indicates that the load is distributed between two serial links. In this scenario, load sharing occurs on a per-destination basis. The **show ip bgp** command gives the valid entries for the 2.0.0.0 network.

```
RouterA# show ip route
!--- Output suppressed. Gateway of last resort is not set 1.0.0.0/24 is subnetted, 1 subnets C
1.1.1.0 is directly connected, Loopback0 B      2.0.0.0/8 [20/0] via 150.10.10.2, 00:04:23
          [20/0] via 160.20.20.2, 00:04:01
160.20.0.0/24 is subnetted, 1 subnets
C      160.20.20.0 is directly connected, Serial0
150.10.0.0/24 is subnetted, 1 subnets
C      150.10.10.0 is directly connected, Serial1
```

```
RouterA# traceroute 2.2.2.2
```

```
Type escape sequence to abort.
Tracing the route to 2.2.2.2
```

```
1 160.20.20.2 16 msec
  150.10.10.2 8 msec *
```

```
RouterA# show ip bgp
```

```
BGP table version is 3, local router ID is 1.1.1.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal
Origin codes: i - IGP, e - EGP, ? - incomplete
```

```
Network Next Hop Metric LocPrf Weight Path
*> 1.0.0.0 0.0.0.0          0 32768 i
*> 2.0.0.0 160.20.20.2    0 0 10 i
*          150.10.10.2    0 0 10 i
```

## Troubleshoot

There is currently no specific troubleshooting information available for this configuration.

# Load Sharing When Dual-Homed to One ISP Through Multiple Local Routers

This scenario shows how to achieve load sharing when there are multiple connections to the same ISP through multiple local routers. The two eBGP peers are terminated on two separate local routers. Load balancing on the two links is not possible because BGP chooses the single best path among the networks that is learned from eBGP and internal BGP (iBGP). Load sharing among the multiple paths to AS 10 is the next-best option. With this type of load sharing, traffic to specific networks, on the basis of predefined policies, travels through both links. Additionally, each link acts as a backup to the other link, in case one link fails.

For simplicity, assume that the BGP routing policy for AS 11 is:


- AS 11 accepts the local routes from AS 10, along with a default for the rest of the Internet routes.
- The outbound traffic policy is: All traffic that is destined to the Internet from R101 goes out through the R101-R103 link. If the R101-R103 link fails, then all traffic to the Internet from R101 goes through R102 to AS 10. Similarly, all traffic that is destined to the Internet from R102 goes through the R102-R104 link. If the R102-R104 link fails, then all traffic to the Internet from R102 goes through R101 to AS 10.
- The inbound traffic policy is: Traffic that is destined for network 192.168.11.0/24 from the Internet should come from the R103-R101 link. Traffic that is destined for network 192.168.12.0/24 from the Internet should come from the R104-R102 link. If one link to AS 10 fails, then the other link should route the traffic that is destined for all networks back to AS 11 from the Internet.

In order to achieve this, 192.168.11.0 is announced from R101 to R103 with a shorter AS\_PATH than is announced from R102 to R104. AS 10 finds the best path through the R103-R101 link. Similarly, 192.168.12.0 is announced with a shorter path through the R102-R104 link. AS 10 prefers the R104-R102 link for traffic that is bound to 192.168.12.0 in AS 11.

For outbound traffic, BGP determines the best path on the basis of routes that are learned through eBGP. These routes are preferable to routes learned through iBGP. So R101 learns 10.10.34.0 from R103 through eBGP and from R102 through iBGP. The external path is selected over the internal path. So, if you look at the BGP table in the [R101](#) configuration, the route toward 10.10.34.0 would be through the R101-R103 link, with next hop 10.10.13.3. On [R102](#), the route toward 10.10.34.0 would be through the R102-R104 link, with next hop 10.10.24.4. This achieves load sharing for traffic that is destined to 10.10.34.0. Similar reasoning applies to the default routes on R101 and R102. For more information on BGP path selection criteria, refer to [BGP Best Path Selection Algorithm](#).

## Network Diagram

This section uses this network setup:



# Configurations

This section uses these configurations:

- [R101](#)
- [R102](#)
- [R103](#)
- [R104](#)

## R101

```
RouterA# show ip route
!--- Output suppressed. Gateway of last resort is not set 1.0.0.0/24 is subnetted, 1 subnets C 1.1.1.0
directly connected, Loopback0 B    2.0.0.0/8 [20/0] via 150.10.10.2, 00:04:23
    [20/0] via 160.20.20.2, 00:04:01
    160.20.0.0/24 is subnetted, 1 subnets
C    160.20.20.0 is directly connected, Serial0
    150.10.0.0/24 is subnetted, 1 subnets
C    150.10.10.0 is directly connected, Serial1

RouterA# traceroute 2.2.2.2

Type escape sequence to abort.
Tracing the route to 2.2.2.2

 1 160.20.20.2 16 msec
   150.10.10.2 8 msec *
```

## R102

```
RouterA# show ip route
!--- Output suppressed. Gateway of last resort is not set 1.0.0.0/24 is subnetted, 1 subnets C 1.1.1.0
directly connected, Loopback0 B    2.0.0.0/8 [20/0] via 150.10.10.2, 00:04:23
    [20/0] via 160.20.20.2, 00:04:01
    160.20.0.0/24 is subnetted, 1 subnets
C    160.20.20.0 is directly connected, Serial0
    150.10.0.0/24 is subnetted, 1 subnets
C    150.10.10.0 is directly connected, Serial1

RouterA# traceroute 2.2.2.2

Type escape sequence to abort.
Tracing the route to 2.2.2.2

 1 160.20.20.2 16 msec
   150.10.10.2 8 msec *
```

```
RouterA# show ip bgp

BGP table version is 3, local router ID is 1.1.1.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal
Origin codes: i - IGP, e - EGP, ? - incomplete

Network Next Hop Metric LocPrf Weight Path
*> 1.0.0.0 0.0.0.0          0 32768 i
*> 2.0.0.0 160.20.20.2    0 0 10 i
*          150.10.10.2    0 0 10 i
```



Origin codes: i - IGP, e - EGP, ? - incomplete

Network Next Hop Metric LocPrf Weight Path

\*> 1.0.0.0 0.0.0.0 0 32768 i

\*> **2.0.0.0 160.20.20.2 0 0 10 i**

\* **150.10.10.2 0 0 10 i**

## R103

RouterA# **show ip route**

*!--- Output suppressed.* Gateway of last resort is not set 1.0.0.0/24 is subnetted, 1 subnets C 1.1.1.0  
directly connected, Loopback0 **B 2.0.0.0/8 [20/0] via 150.10.10.2, 00:04:23**  
**[20/0] via 160.20.20.2, 00:04:01**

160.20.0.0/24 is subnetted, 1 subnets

C 160.20.20.0 is directly connected, Serial0

150.10.0.0/24 is subnetted, 1 subnets

C 150.10.10.0 is directly connected, Serial1

RouterA# **traceroute 2.2.2.2**

Type escape sequence to abort.

Tracing the route to 2.2.2.2

**1 160.20.20.2 16 msec**  
**150.10.10.2 8 msec \***

RouterA# **show ip bgp**

BGP table version is 3, local router ID is 1.1.1.1

Status codes: s suppressed, d damped, h history, \* valid, > best, i - internal

Origin codes: i - IGP, e - EGP, ? - incomplete

Network Next Hop Metric LocPrf Weight Path

\*> 1.0.0.0 0.0.0.0 0 32768 i

\*> **2.0.0.0 160.20.20.2 0 0 10 i**

\* **150.10.10.2 0 0 10 i**

## R104

RouterA# **show ip route**

*!--- Output suppressed.* Gateway of last resort is not set 1.0.0.0/24 is subnetted, 1 subnets C 1.1.1.0  
directly connected, Loopback0 **B 2.0.0.0/8 [20/0] via 150.10.10.2, 00:04:23**  
**[20/0] via 160.20.20.2, 00:04:01**

160.20.0.0/24 is subnetted, 1 subnets

C 160.20.20.0 is directly connected, Serial0

150.10.0.0/24 is subnetted, 1 subnets

C 150.10.10.0 is directly connected, Serial1

RouterA# **traceroute 2.2.2.2**

Type escape sequence to abort.

Tracing the route to 2.2.2.2

**1 160.20.20.2 16 msec**  
**150.10.10.2 8 msec \***

RouterA# **show ip bgp**

BGP table version is 3, local router ID is 1.1.1.1

Status codes: s suppressed, d damped, h history, \* valid, > best, i - internal

Origin codes: i - IGP, e - EGP, ? - incomplete

Network Next Hop Metric LocPrf Weight Path

\*> 1.0.0.0 0.0.0.0 0 32768 i

```
*> 2.0.0.0 160.20.20.2 0 0 10 i
*
      150.10.10.2 0 0 10 i
```

## Verify

This section provides information you can use to confirm your configuration is working properly.

Certain **show** commands are supported by the [Cisco CLI Analyzer](#) ([registered](#) customers only) , which allows you to view an analysis of **show** command output.

## Verification When Both Links Between AS 11 and AS 10 Are Up

### Verification of Outbound Traffic

**Note:** The greater-than sign (>) in the **show ip bgp** command output represents the best path to use for that network among the possible paths. Refer to [BGP Best Path Selection Algorithm](#) for more information.

The BGP table in [R101](#) shows that the best path for all outgoing traffic to the Internet is through the R101-R103 link. The **show ip route** command output confirms the routes in the routing table.

```
R101# show ip bgp
```

```
BGP table version is 5, local router ID is 192.168.12.1
```

```
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal
```

```
Origin codes: i - IGP, e - EGP, ? - incomplete
```

| Network          | Next Hop                                  | Metric | LocPrf | Weight   | Path   |
|------------------|---|--------|--------|----------|--|
| * i0.0.0.0       | 192.168.12.2                              |        | 100    | 0        | 10 i   |
| <b>*&gt;</b>     | <b>10.10.13.3</b>                         |        |        | <b>0</b> | <b>10 i</b> <i>!--- This is the next hop of R103.</i>                        |
| * i10.10.34.0/24 | 192.168.12.2                              | 100    | 0      | 10 i     | <b>*&gt;</b> <b>10.10.13.3</b> <b>0</b>                                      |
| <b>0 10 i</b>    | <i>!--- This is the next hop of R103.</i> |        |        |          | <b>* i192.168.11.0</b> 192.168.12.2 0 100 0 i <b>*&gt;</b> 0.0.0.0 0 32768 i |
| * i192.168.12.0  | 192.168.12.2                              | 0      | 100    | 0        | i <b>*&gt;</b> 0.0.0.0 0 32768 i   |

```
R101# show ip route
!--- Output suppressed. Gateway of last resort is 10.10.13.3 to network 0.0.0.0 C
192.168.12.0/24 is directly connected, Ethernet0/0 C 192.168.11.0/24 is directly connected,
Ethernet0/0 10.0.0.0/24 is subnetted, 2 subnets C 10.10.13.0 is directly connected, Serial8/0 B
10.10.34.0 [20/0] via 10.10.13.3, 00:08:53 !--- This is the next hop of R103. B* 0.0.0.0/0
[20/0] via 10.10.13.3, 00:08:53 !--- This is the next hop of R103.
```

Here are the BGP and routing tables for R102. According to the policy, R102 should route all the traffic to AS 10 through the R102-R104 link:

```
R102# show ip bgp
```

```
BGP table version is 7, local router ID is 192.168.12.2
```

```
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal
```

```
Origin codes: i - IGP, e - EGP, ? - incomplete
```

| Network              | Next Hop                                  | Metric | LocPrf | Weight   | Path  |
|----------------------|---|--------|--------|--|---|
| <b>*&gt;</b> 0.0.0.0 | <b>10.10.24.4</b>                         |        |        | <b>0</b>   | <b>10 i</b> <i>!--- This is the next hop of R104.</i>                         |
| * i 192.168.12.1     | 100                                       | 0      | 10 i   | <b>*&gt;</b> <b>10.10.34.0/24</b> <b>10.10.24.4</b> <b>0</b> |   |
| <b>10 i</b>          | <i>!--- This is the next hop of R104.</i> |        |        |  | <b>* i 192.168.11.0</b> 192.168.12.1 0 100 0 i <b>*&gt;</b> 0.0.0.0 0 32768 i |
| * i 192.168.12.0     | 192.168.12.1                              | 0      | 100    | 0  | i <b>*&gt;</b> 0.0.0.0 0 32768 i  |

```
R102# show ip route
!--- Output suppressed. Gateway of last resort is 10.10.24.4 to network 0.0.0.0 C
192.168.12.0/24 is directly connected, Ethernet0/0 C 192.168.11.0/24 is directly connected,
Ethernet0/0 10.0.0.0/24 is subnetted, 2 subnets C 10.10.24.0 is directly connected, Serial8/0 B
10.10.34.0 [20/0] via 10.10.24.4, 00:11:21 !--- This is the next hop of R104. B* 0.0.0.0/0
[20/0] via 10.10.24.4, 00:11:21 !--- This is the next hop of R104.
```

## Verification of Inbound Traffic from AS 10 to AS 11

The networks 192.168.11.0 and 192.168.12.0 belong to AS 11. According to the policy, AS 11 should prefer the R103-R101 link for traffic that is destined to network 192.168.11.0 and the R104-R102 link for traffic that is destined to network 192.168.12.0.

```
R103# show ip bgp
```

BGP table version is 4, local router ID is 10.10.34.3

Status codes: s suppressed, d damped, h history, \* valid, > best, i - internal

Origin codes: i - IGP, e - EGP, ? - incomplete

| Network          | Next Hop   | Metric            | LocPrf | Weight | Path                       |
|------------------|------------|-------------------|--------|--------|----------------------------|
| *> 10.10.34.0/24 | 0.0.0.0    | 0                 |        | 32768  | i                          |
| *> 192.168.11.0  | 10.10.13.1 | 0                 |        | 0 11 i | !--- The next hop is R101. |
| * 192.168.12.0   | 10.10.13.1 | 0 0 11 11 11 11 i |        |        | *>i 10.10.34.4             |
| 0 100            | 0 11 i     |                   |        |        | !--- The next hop is R104. |

R103# show ip route

!--- Output suppressed. Gateway of last resort is not set B 192.168.12.0/24 [200/0] via 10.10.34.4, 00:04:46 !--- The next hop is R104. B 192.168.11.0/24 [20/0] via 10.10.13.1, 00:04:46 !--- The next hop is R101. 10.0.0.0/24 is subnetted, 2 subnets C 10.10.13.0 is directly connected, Serial8/0 C 10.10.34.0 is directly connected, Ethernet0/0

The best path for network 192.168.11.0 on R103 is through the R103-R101 link, and the best path for network 192.168.12.0 is through R104 to AS 11. In this case, the shortest path length determines the best path.

Similarly, on R104, the BGP and routing table resembles this:

```
R104# show ip bgp
```

BGP table version is 13, local router ID is 10.10.34.4

Status codes: s suppressed, d damped, h history, \* valid, > best, i - internal

Origin codes: i - IGP, e - EGP, ? - incomplete


| Network          | Next Hop   | Metric | LocPrf | Weight          | Path |
|------------------|------------|--------|--------|-----------------|------|
| *>i10.10.34.0/24 | 10.10.34.3 | 0      | 100    | 0               | i    |
| *>i192.168.11.0  | 10.10.34.3 | 0      | 100    | 0 11 i          |      |
| *                | 10.10.24.2 | 0      |        | 0 11 11 11 11 i |      |
| *> 192.168.12.0  | 10.10.24.2 | 0      |        | 0 11 i          |      |

```
R104# show ip route
```

!--- Output suppressed. Gateway of last resort is not set B 192.168.12.0/24 [20/0] via 10.10.24.2, 00:49:06 !--- The next hop is R102. B 192.168.11.0/24 [200/0] via 10.10.34.3, 00:07:36 !--- The next hop is R103. 10.0.0.0/24 is subnetted, 2 subnets C 10.10.24.0 is directly connected, Serial8/0 C 10.10.34.0 is directly connected, Ethernet0/0

## Verification When the R101-R103 Link Fails

When the R101-R103 link fails, all traffic should reroute through R102. This diagram illustrates this change:



Shut down the R103-R101 link on R103 in order to simulate this situation.

```
R103(config)# interface serial 8/0
R103(config-if)# shutdown
```

```
*May 1 00:52:33.379: %BGP-5-ADJCHANGE: neighbor 10.10.13.1 Down Interface flap
*May 1 00:52:35.311: %LINK-5-CHANGED: Interface Serial8/0, changed state to
administratively down
*May 1 00:52:36.127: %LINEPROTO-5-UPDOWN: Line protocol on Interface Serial8/0, changed
state to down
```

Verify the outbound route to AS 10.

```
R101# show ip bgp
```

```
BGP table version is 17, local router ID is 192.168.12.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal
Origin codes: i - IGP, e - EGP, ? - incomplete
```

| Network   | Next Hop     | Metric | LocPrf | Weight | Path |
|---|--------------|--------|--------|--------|------|
| *>i0.0.0.0  | 192.168.12.2 | 100    | 0      | 10     | i    |
| <i>!--- This is the next hop of R102.</i>   |              |        |        |        |      |
| *>i10.10.34.0/24  | 192.168.12.2 | 100    | 0      | 10     | i    |
| <i>!--- This is the next hop of R102.</i>   |              |        |        |        |      |
| * i192.168.11.0 192.168.12.2 0 100 0 i *> 0.0.0.0 0 32768 i * i192.168.12.0 192.168.12.2 0 100 0 i *> 0.0.0.0 0 32768 i                       |              |        |        |        |      |
| R101# show ip route   |              |        |        |        |      |
| <i>!--- Output suppressed.</i> Gateway of last resort is 192.168.12.2 to network 0.0.0.0 C  |              |        |        |        |      |
| 192.168.12.0/24 is directly connected, Ethernet0/0 C 192.168.11.0/24 is directly connected, Ethernet0/0 10.0.0.0/24 is subnetted, 1 subnets B |              |        |        |        |      |
| 10.10.34.0 [200/0] via 192.168.12.2, 00:01:34   |              |        |        |        |      |
| B* 0.0.0.0/0 [200/0] via 192.168.12.2, 00:01:34   |              |        |        |        |      |
| <i>!--- All outbound traffic goes through R102.</i>   |              |        |        |        |      |
| R102# show ip route   |              |        |        |        |      |
| <i>!--- Output suppressed.</i> Gateway of last resort is 10.10.24.4 to network 0.0.0.0 C  |              |        |        |        |      |
| 192.168.12.0/24 is directly connected, Ethernet0/0 C 192.168.11.0/24 is directly connected, Ethernet0/0 10.0.0.0/24 is subnetted, 2 subnets C |              |        |        |        |      |
| 10.10.24.0 is directly connected, Serial8/0 B   |              |        |        |        |      |
| 10.10.34.0 [20/0] via 10.10.24.4, 00:13:22  |              |        |        |        |      |
| B* 0.0.0.0/0 [20/0] via 10.10.24.4, 00:55:22  |              |        |        |        |      |
| <i>!--- All outbound traffic on R102 goes through R104.</i>   |              |        |        |        |      |

Verify the inbound traffic route when R101-R103 is down.

```
R103# show ip bgp
```

```
BGP table version is 6, local router ID is 10.10.34.3
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal
Origin codes: i - IGP, e - EGP, ? - incomplete
```

| Network          | Next Hop   | Metric | LocPrf | Weight | Path          |
|------------------|------------|--------|--------|--------|---------------|
| *> 10.10.34.0/24 | 0.0.0.0    | 0      |        | 32768  | i             |
| *>i192.168.11.0  | 10.10.34.4 | 0      | 100    | 0      | 11 11 11 11 i |
| *>i192.168.12.0  | 10.10.34.4 | 0      | 100    | 0      | 11 i          |

```
R103# show ip route
```

```
!--- Output suppressed. Gateway of last resort is not set B 192.168.12.0/24 [200/0] via
10.10.34.4, 00:14:55 !--- The next hop is R104. B 192.168.11.0/24 [200/0] via
10.10.34.4, 00:05:46 !--- The next hop is R104. 10.0.0.0/24 is subnetted, 1 subnets C
10.10.34.0 is directly connected, Ethernet0/0
```

On R104, the traffic for 192.168.11.0 and 192.168.12.0 goes through the R104-R102 link.

```
R104# show ip route
```

```
!--- Output suppressed. Gateway of last resort is not set B 192.168.12.0/24 [20/0] via
10.10.24.2, 00:58:35 !--- The next hop is R102. B 192.168.11.0/24 [20/0] via 10.10.24.2,
00:07:57 !--- The next hop is R102. 10.0.0.0/24 is subnetted, 2 subnets C 10.10.24.0 is
directly connected, Serial8/0 C 10.10.34.0 is directly connected, Ethernet0/0
```

## Troubleshoot

There is currently no specific troubleshooting information available for this configuration.


## Load Sharing When Multihomed to Two ISPs Through a Single Local Router

In this scenario, load balancing is not an option in a multihomed environment, so you can only do load sharing. You cannot do load balancing because BGP selects only a single best path to a destination among the BGP routes that are learned from the different ASs. The idea is to set a better metric for the routes in the range 1.0.0.0 to 128.0.0.0 that are learned from ISP(A) and a better metric for the rest of the routes that are learned from ISP(B). The [Network Diagram](#) is an example.

Refer to [Sample Configuration for BGP with Two Different Service Providers \(Multihoming\)](#) for additional information.

### Network Diagram

This section uses this network setup:



### Configurations

This section uses these configurations:

- [RouterA](#)
- [RouterB](#)
- [RouterC](#)

#### RouterA

```
R104# show ip route
```

```
!--- Output suppressed. Gateway of last resort is not set B 192.168.12.0/24 [20/0] via 10.10.24.2, 00:58:35
!--- The next hop is R102. B 192.168.11.0/24 [20/0] via 10.10.24.2, 00:07:57
!--- The hop is R102. 10.0.0.0/24 is subnetted, 2 subnets C 10.10.24.0 is directly connected, Serial8/0 C 10.10. is directly connected, Ethernet0/0
```

## RouterB

```
R104# show ip route
```

```
!--- Output suppressed. Gateway of last resort is not set B    192.168.12.0/24 [20/0] via 10.10.24.2,
00:58:35    !--- The next hop is R102. B    192.168.11.0/24 [20/0] via 10.10.24.2, 00:07:57    !--- The
hop is R102. 10.0.0.0/24 is subnetted, 2 subnets C 10.10.24.0 is directly connected, Serial8/0 C 10.10.
is directly connected, Ethernet0/0
```

## RouterC

```
R104# show ip route
```

```
!--- Output suppressed. Gateway of last resort is not set B    192.168.12.0/24 [20/0] via 10.10.24.2,
00:58:35    !--- The next hop is R102. B    192.168.11.0/24 [20/0] via 10.10.24.2, 00:07:57    !--- The
hop is R102. 10.0.0.0/24 is subnetted, 2 subnets C 10.10.24.0 is directly connected, Serial8/0 C 10.10.
is directly connected, Ethernet0/0
```

## Verify

Use this section to confirm that your configuration works properly.

The [Cisco CLI Analyzer](#) ([registered](#) customers only) supports certain **show** commands. Use the Cisco CLI Analyzer to view an analysis of **show** command output.

The output from the **show ip route** command and the output from the **traceroute** command show that any network lower than 128.0.0.0 exits RouterA through 160.20.20.2. This route is the next hop out of the serial 0 interface. The rest of the networks exit through 150.10.10.2, which is the next hop out of the serial 1 interface.

```
RouterA# show ip route
```

```
!--- Output suppressed. Gateway of last resort is not set B 170.16.0.0/16 [20/0] via
150.10.10.2, 00:43:43 !--- This is the next hop out through serial 1. B 2.0.0.0/8 [20/0] via
160.20.20.2, 00:43:43 !--- This is the next hop out through serial 0. 160.20.0.0/24 is
subnetted, 1 subnets C 160.20.20.0 is directly connected, Serial0 150.10.0.0/24 is subnetted, 1
subnets C 150.10.10.0 is directly connected, Serial1 RouterA# show ip bgp
```

```
BGP table version is 3, local router ID is 160.20.20.1
```

```
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal
```

```
Origin codes: i - IGP, e - EGP, ? - incomplete
```

| Network      | Next Hop    | Metric | LocPrf | Weight | Path |
|--------------|-------------|--------|--------|--------|------|
| * 2.0.0.0    | 150.10.10.2 | 0      |        | 0 12   | i    |
| *>           | 160.20.20.2 | 0      |        | 100 10 | i    |
| * 170.16.0.0 | 160.20.20.2 | 0      |        | 0 10   | i    |
| *>           | 150.10.10.2 | 0      |        | 100 12 | i    |

```
RouterA# traceroute 2.2.2.2
```

```
Type escape sequence to abort.
```

```
Tracing the route to 2.2.2.2
```

```
1 160.20.20.2 16 msec * 16 msec
```

```
RouterA# traceroute 170.16.6.6
```

```
Type escape sequence to abort.
```

```
Tracing the route to 170.16.6.6
```

```
1 150.10.10.2 4 msec * 4 msec
```

## Troubleshoot

There is currently no specific troubleshooting information available for this configuration.

## Load Sharing When Multihomed to Two ISPs Through Multiple Local Routers

Load balancing is not possible in a multihomed environment with two ISPs. BGP selects only the single best path to a destination among the BGP paths that are learned from different ASs, which makes load balancing impossible. But, load sharing is possible in such multihomed BGP networks. On the basis of predetermined policies, traffic flow is controlled with different BGP attributes.

This section discusses the multihoming configuration that has the most frequent use. The configuration shows how to achieve load sharing. See the [Network Diagram](#), in which the multihome of AS 100 achieves reliability and load sharing.


**Note:** The IP addresses in this example follow [RFC 1918](#) standards for Private Address Space and are not routable on the Internet.

For simplicity, assume that the BGP routing policy for AS 100 is:

- AS 100 accepts the local routes from both providers, along with a default for the rest of the Internet routes.
- The outbound traffic policy is: Traffic that is destined to AS 300 goes through the R1-ISP(A) link. Traffic that is destined to AS 400 goes through the R2-ISP(B) link. All other traffic should prefer default route 0.0.0.0 through the R1-ISP(A) link. If the R1-ISP(A) link fails, all traffic should go through the R2-ISP(B) link.
- The inbound traffic policy is: Traffic that is destined for network 10.10.10.0/24 from the Internet should come from the ISP(A)-R1 link. Traffic that is destined for network 10.10.20.0/24 from the Internet should come from the ISP(B)-R2 link. If one ISP fails, the other ISP should route traffic back to AS 100 from the Internet for all the networks.

### Network Diagram

This section uses this network setup:



# Configurations

This section uses these configurations:

- [R2](#)
- [R1](#)

## R2

```
RouterA# show ip route
!--- Output suppressed. Gateway of last resort is not set B 170.16.0.0/16 [20/0] via 150.10.10.2, 00:43:43
-- This is the next hop out through serial 1. B 2.0.0.0/8 [20/0] via 160.20.20.2, 00:43:43 !--- This is the next hop out through serial 0. 160.20.0.0/24 is subnetted, 1 subnets C 160.20.20.0 is directly connected
Serial0 150.10.0.0/24 is subnetted, 1 subnets C 150.10.10.0 is directly connected, Serial1 RouterA# show ip route bgp
```

```
BGP table version is 3, local router ID is 160.20.20.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal
Origin codes: i - IGP, e - EGP, ? - incomplete
Network Next Hop Metric LocPrf Weight Path
```

|    | Network    | Next Hop    | Metric | LocPrf | Weight | Path |
|----|------------|-------------|--------|--------|--------|------|
| *  | 2.0.0.0    | 150.10.10.2 | 0      |        | 0 12   | i    |
| *> |            | 160.20.20.2 | 0      |        | 100 10 | i    |
| *  | 170.16.0.0 | 160.20.20.2 | 0      |        | 0 10   | i    |
| *> |            | 150.10.10.2 | 0      |        | 100 12 | i    |

```
RouterA# traceroute 2.2.2.2
```

```
Type escape sequence to abort.
Tracing the route to 2.2.2.2
```

```
1 160.20.20.2 16 msec * 16 msec
```

```
RouterA# traceroute 170.16.6.6
```

```
Type escape sequence to abort.
Tracing the route to 170.16.6.6
```

```
1 150.10.10.2 4 msec * 4 msec
```

## R1

```
RouterA# show ip route
!--- Output suppressed. Gateway of last resort is not set B 170.16.0.0/16 [20/0] via 150.10.10.2, 00:43:43
-- This is the next hop out through serial 1. B 2.0.0.0/8 [20/0] via 160.20.20.2, 00:43:43 !--- This is the next hop out through serial 0. 160.20.0.0/24 is subnetted, 1 subnets C 160.20.20.0 is directly connected
Serial0 150.10.0.0/24 is subnetted, 1 subnets C 150.10.10.0 is directly connected, Serial1 RouterA# show ip route bgp
```

```
BGP table version is 3, local router ID is 160.20.20.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal
Origin codes: i - IGP, e - EGP, ? - incomplete
Network Next Hop Metric LocPrf Weight Path
```

|    | Network    | Next Hop    | Metric | LocPrf | Weight | Path |
|----|------------|-------------|--------|--------|--------|------|
| *  | 2.0.0.0    | 150.10.10.2 | 0      |        | 0 12   | i    |
| *> |            | 160.20.20.2 | 0      |        | 100 10 | i    |
| *  | 170.16.0.0 | 160.20.20.2 | 0      |        | 0 10   | i    |
| *> |            | 150.10.10.2 | 0      |        | 100 12 | i    |

```
RouterA# traceroute 2.2.2.2
```

```
Type escape sequence to abort.
```



Tracing the route to 2.2.2.2

```
1 160.20.20.2 16 msec * 16 msec
```

```
RouterA# traceroute 170.16.6.6
```

Type escape sequence to abort.

Tracing the route to 170.16.6.6

```
1 150.10.10.2 4 msec * 4 msec
```

## Verify

Use this section to confirm that your configuration works properly.

The [Cisco CLI Analyzer](#) ([registered](#) customers only) supports certain **show** commands. Use the Cisco CLI Analyzer to view an analysis of **show** command output.

Issue the **show ip bgp** command in order to verify that the outbound/inbound policy works.

**Note:** The greater-than sign (>) in the **show ip bgp** output represents the best path to use for that network among the possible paths. Refer to [BGP Best Path Selection Algorithm](#) for more information.

```
R1# show ip bgp
```

BGP table version is 6, local router ID is 192.168.31.1

Status codes: s suppressed, d damped, h history, \* valid, > best, i - internal

Origin codes: i - IGP, e - EGP, ? - incomplete

BGP table version is 6, local router ID is 192.168.31.1

Status codes: s suppressed, d damped, h history, \* valid, > best, i - internal

Origin codes: i - IGP, e - EGP, ? - incomplete

| Network  | Next Hop     | Metric | LocPrf | Weight  | Path    |
|--|--------------|--------|--------|---------|---------|
| *> 0.0.0.0   | 192.168.31.3 |        |        | 200     | 0 300 i |
| <i>!--- This line shows that the default route 0.0.0.0/0 is preferred</i>  |              |        |        |         |         |
| <i>!--- through AS 300, ISP(A). * i10.10.10.0/24 192.168.21.2 0 100 0 i *&gt; 0.0.0.0 0 32768 i * i10.10.20.0/24 192.168.21.2 0 100 0 i *&gt; 0.0.0.0 0 32768 i *&gt; 30.30.30.0/24 192.168.31.3</i> |              |        |        |         |         |
| 0  | 200          |        |        | 0 300 i |         |
| *>i40.40.40.0/24   | 192.168.21.2 |        |        | 0 150   | 0 400 i |
| <i>!--- The route to network 30.30.30.0/24 (AS 300) is preferred</i>   |              |        |        |         |         |
| <i>!--- through the R1-ISP(A) link.</i>  |              |        |        |         |         |
| <i>!--- The route to network 40.40.40.0/24 (AS 400) is preferred</i>   |              |        |        |         |         |
| <i>!--- through the R2-ISP(B) link.</i>  |              |        |        |         |         |

Now, look at the **show ip bgp** output on R2:

```
R2# show ip bgp
```

BGP table version is 8, local router ID is 192.168.42.2

Status codes: s suppressed, d damped, h history, \* valid, > best, i - internal

Origin codes: i - IGP, e - EGP, ? - incomplete

| Network   | Next Hop     | Metric | LocPrf | Weight | Path    |
|---|--------------|--------|--------|--------|---------|
| * 0.0.0.0   | 192.168.42.4 |        |        | 150    | 0 400 i |
| *>i   | 192.168.21.1 |        |        | 200    | 0 300 i |
| <i>!--- This line shows that the default route 0.0.0.0/0 is preferred</i>   |              |        |        |        |         |
| <i>!--- through AS 300, through the R2-ISP(B) link. *&gt; 10.10.10.0/24 0.0.0.0 0 32768 i * i 192.168.21.1 0 100 0 i *&gt; 10.10.20.0/24 0.0.0.0 0 32768 i * i 192.168.21.1 0 100 0 i</i> |              |        |        |        |         |
| *>i30.30.30.0/24  | 192.168.21.1 |        |        | 0 200  | 0 300 i |

```
*> 40.40.40.0/24      192.168.42.4          0      150      0 400 i
!--- The route to network 30.30.30.0/24 (AS 300) is preferred
!--- through the R1-ISP(A) link.
!--- The route to network 40.40.40.0/24 (AS 400) is preferred
!--- through the R2-ISP(B) link.
```

Issue the **show ip bgp** command on Router 6 in order to observe the inbound policy for networks 10.10.10.0/24 and 10.10.20.0/24:

```
R6# show ip bgp
```

```
BGP table version is 15, local router ID is 192.168.64.6
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal
Origin codes: i - IGP, e - EGP, ? - incomplete
```

| Network  | Next Hop     | Metric | LocPrf | Weight        | Path |
|--|--------------|--------|--------|---------------|------|
| *> 10.10.10.0/24   | 192.168.63.3 |        |        | 0 300 100 100 | i    |
| !--- This line shows that network 10.10.10.0/24 is routed through AS 300   |              |        |        |               |      |
| !--- with the ISP(A)-R1 link. * 192.168.64.4 0 400 100 100 100 i * 10.10.20.0/24 192.168.63.3 0 300 100 100 i *> |              |        |        |               |      |
|  | 192.168.64.4 |        |        | 0 400 100     | i    |
| !--- This line shows that network 10.10.20.0/24 is routed through AS 400   |              |        |        |               |      |
| !--- with the ISP(B)-R2 link. *> 30.30.30.0/24 192.168.63.3 0 0 300 i *> 40.40.40.0/24 192.168.64.4 0 0 400 i    |              |        |        |               |      |

Shut down the R1-ISP(A) link on R1 and observe the BGP table. Expect all traffic to the Internet to be routed through the R2-ISP(B) link:

```
R1(config)# interface serial 0/0
R1(config-if)# shutdown
```

```
*May  2 19:00:47.377: %BGP-5-ADJCHANGE: neighbor 192.168.31.3 Down Interface flap
*May  2 19:00:48.277: %LINK-5-CHANGED: Interface Serial0/0, changed state to
administratively down
*May 23 12:00:51.255: %LINEPROTO-5-UPDOWN: Line protocol on Interface Serial0, changed
state to down
```

```
R1# show ip bgp
```

```
BGP table version is 12, local router ID is 192.168.31.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal
Origin codes: i - IGP, e - EGP, ? - incomplete
```

| Network   | Next Hop     | Metric | LocPrf | Weight | Path    |
|---|--------------|--------|--------|--------|---------|
| *>i0.0.0.0  | 192.168.21.2 |        |        | 150    | 0 400 i |
| !--- The best default path is now through the R2-ISP(B) link. * i10.10.10.0/24 192.168.21.2 0 100 0 i *> 0.0.0.0 0 32768 i * i10.10.20.0/24 192.168.21.2 0 100 0 i *> 0.0.0.0 0 32768 i |              |        |        |        |         |
| *>i40.40.40.0/24 192.168.21.2 0 150 0 400 i R2# show ip bgp   |              |        |        |        |         |

```
BGP table version is 14, local router ID is 192.168.42.2
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal
Origin codes: i - IGP, e - EGP, ? - incomplete
```

| Network   | Next Hop     | Metric | LocPrf | Weight | Path    |
|---|--------------|--------|--------|--------|---------|
| *> 0.0.0.0  | 192.168.42.4 |        |        | 150    | 0 400 i |
| !--- The best default route is now through ISP(B) with a  |              |        |        |        |         |
| !--- local preference of 150. * i10.10.10.0/24 192.168.21.1 0 100 0 i *> 0.0.0.0 0 32768 i * i10.10.20.0/24 192.168.21.1 0 100 0 i *> 0.0.0.0 0 32768 i *> 40.40.40.0/24 192.168.42.4 0 150 0 400 i |              |        |        |        |         |

Look at the route for network 10.10.10.0/24 in Router 6:

```
R6# show ip bgp
```

```
BGP table version is 14, local router ID is 192.168.64.6
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal
```

Origin codes: i - IGP, e - EGP, ? - incomplete

```

      Network          Next Hop          Metric LocPrf Weight Path
*> 10.10.10.0/24      192.168.64.4                      0 400 100 100 i
  !--- Network 10.10.10.0 is reachable through ISP(B), which announced
  !--- the network with AS path prepend.  *> 10.10.20.0/24 192.168.64.4 0 400 100 i  *>
30.30.30.0/24 192.168.63.3 0 0 300 i  *> 40.40.40.0/24 192.168.64.4 0 0 400 i
```

## Troubleshoot

There is currently no specific troubleshooting information available for this configuration.

## Related Information

- [BGP Multi-homing: Design and Troubleshooting - Video from live Webcast](#)
- [BGP Multi-homing: Design and Troubleshooting - Questions and Answers from live Webcast](#)
- [How Does Load Balancing Work?](#)
- [Sample Configuration for BGP with Two Different Service Providers \(Multihoming\)](#)
- [How BGP Routers Use the Multi-Exit Discriminator for Best Path Selection](#)
- [BGP Technology Support Page](#)
- [IP Routing Technology Support Page](#)
- [Technical Support & Documentation - Cisco Systems](#)